



CFG/SSA/LLVM Notes

Lecturer: Sanjay Madhav

Control Flow Graph

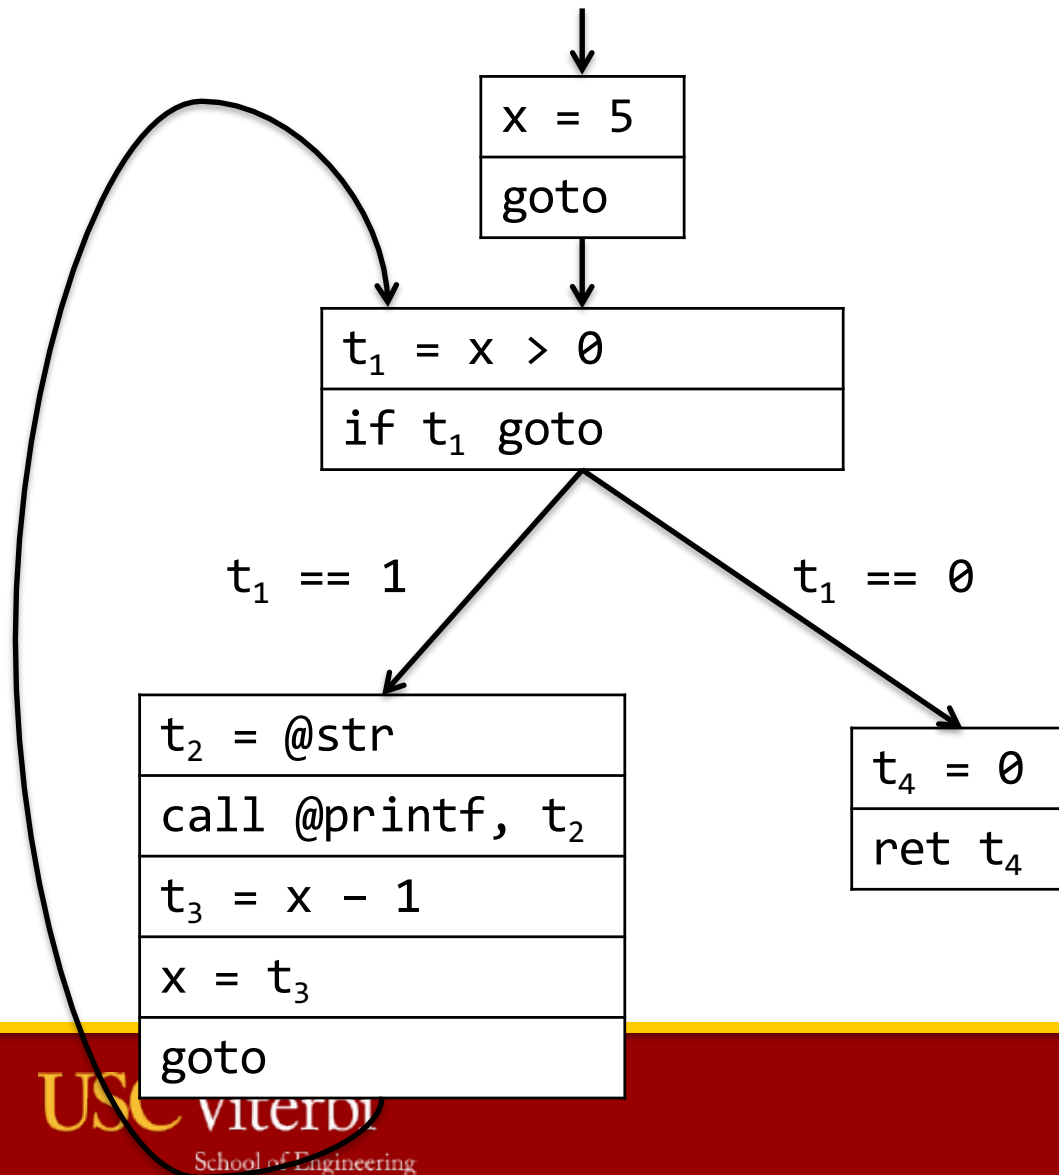


- A **control flow graph** (CFG) is a hybrid IR
- In a CFG, sequences of linear code (without any jumps, gotos, branches, or the like) are called **basic blocks**
- Basic blocks are usually represented by 3 address code (or similar)
- Jumps in control flow are represented by edges in the graph

Control Flow Graph, Cont'd



- Example of a CFG

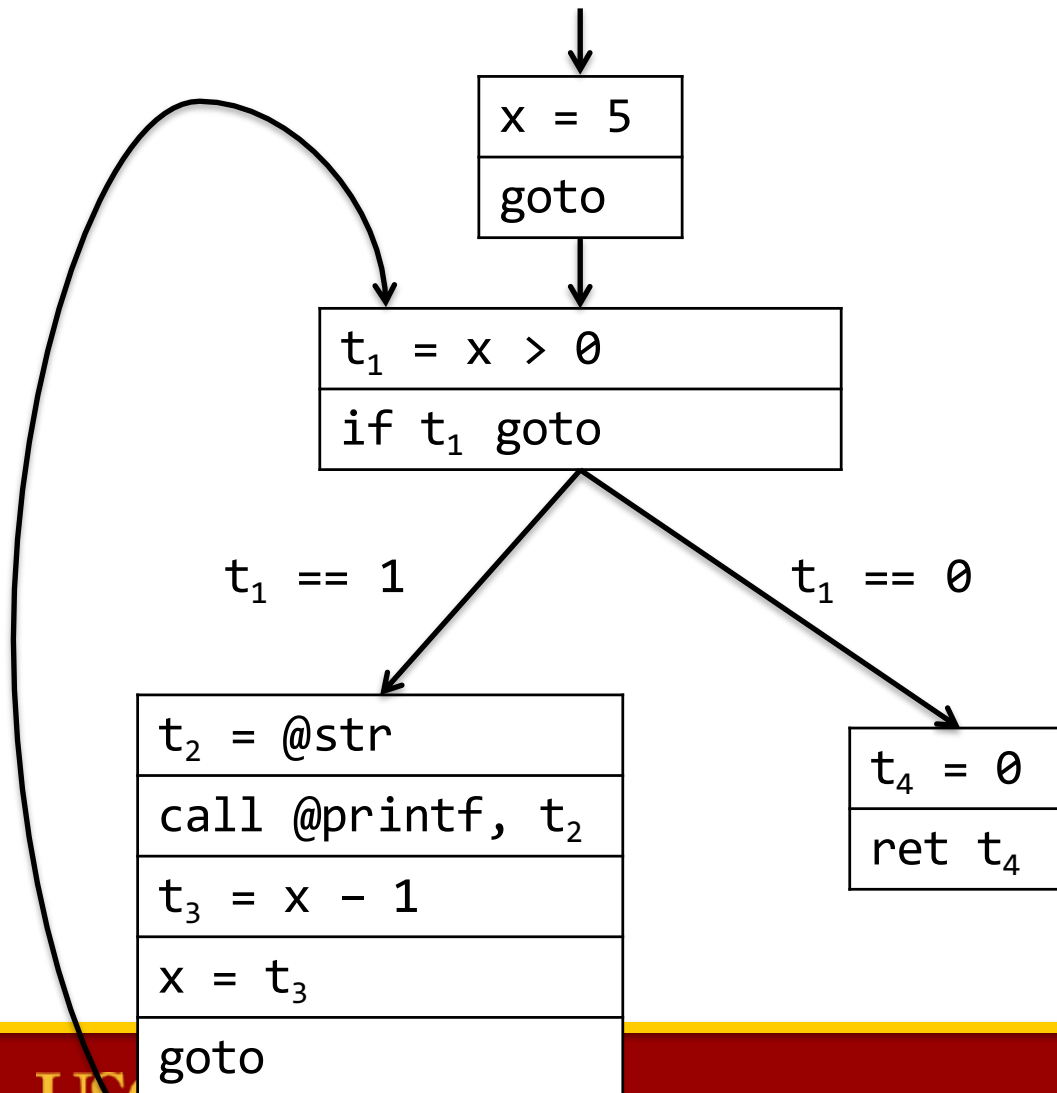


```
int main() {  
    int x = 5;  
    while (x > 0) {  
        printf("Hello!");  
        --x;  
    }  
    return 0;  
}
```

Comparing CFG to 3 Address Code



- Because of the “goto 2”, the “x = 5” must be in a separate block

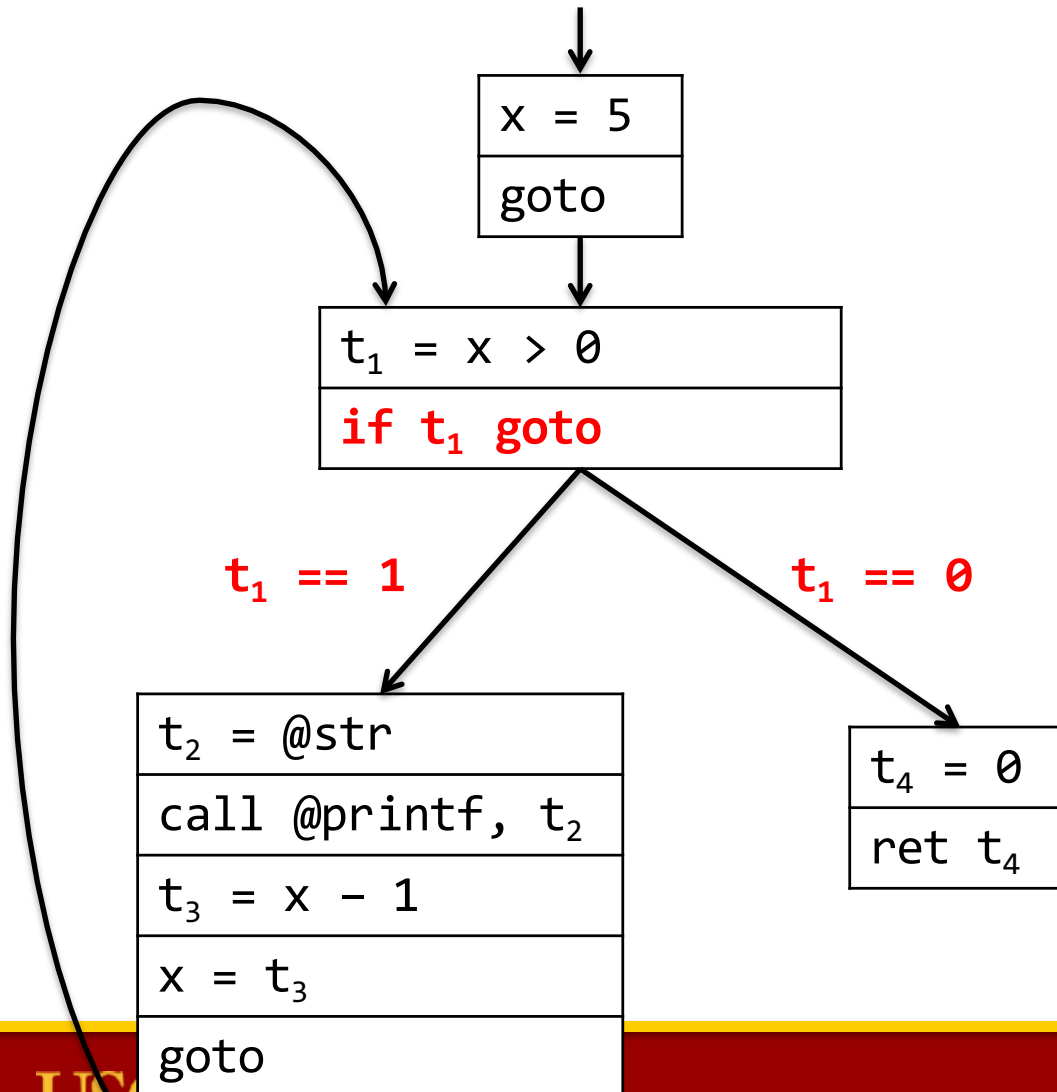


1:	x = 5
2:	t ₁ = x > 0
3:	if !t ₁ goto 10
4:	t ₂ = @str
5:	t ₃ = @printf
6:	call t ₃ , t ₂
7:	t ₄ = x - 1
8:	x = t ₄
9:	goto 2
10:	t ₅ = 0
11:	ret t ₅

Comparing CFG to 3 Address Code



- The way I wrote the if/goto is slightly different, also

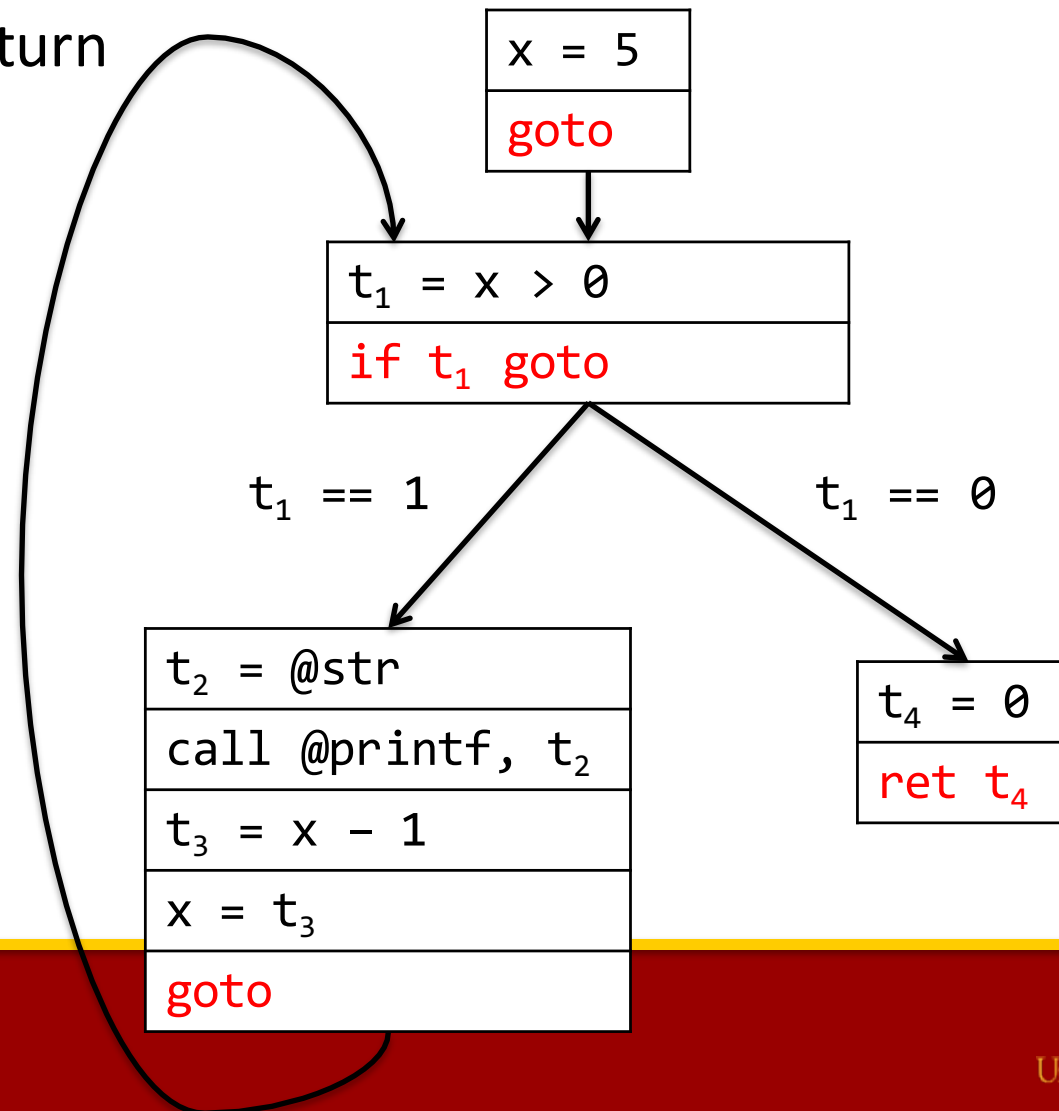


1:	x = 5
2:	t ₁ = x > 0
3:	if !t ₁ goto 10
4:	t ₂ = @str
5:	t ₃ = @printf
6:	call t ₃ , t ₂
7:	t ₄ = x - 1
8:	x = t ₄
9:	goto 2
10:	t ₅ = 0
11:	ret t ₅



Control Flow Graph, Cont'd

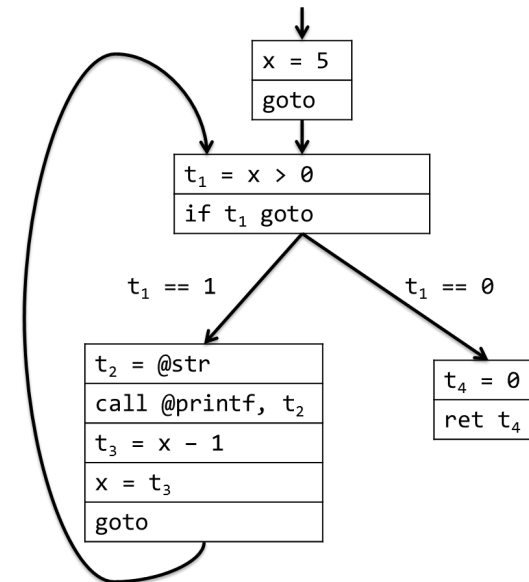
- Each basic block *must* end with a **terminator instruction**
- The terminator is either a branch (conditional or unconditional) or a function return



CFG Advantages/Disadvantages



- Advantages:
 - Clearly represents control flow, which allows for loops to be optimized
 - Each basic block is guaranteed to have sequential code and so it can be aggressively optimized
- Disadvantages:
 - Cannot be generated at parse time
 - Requires the most amount of code to create, out of all the IR covered

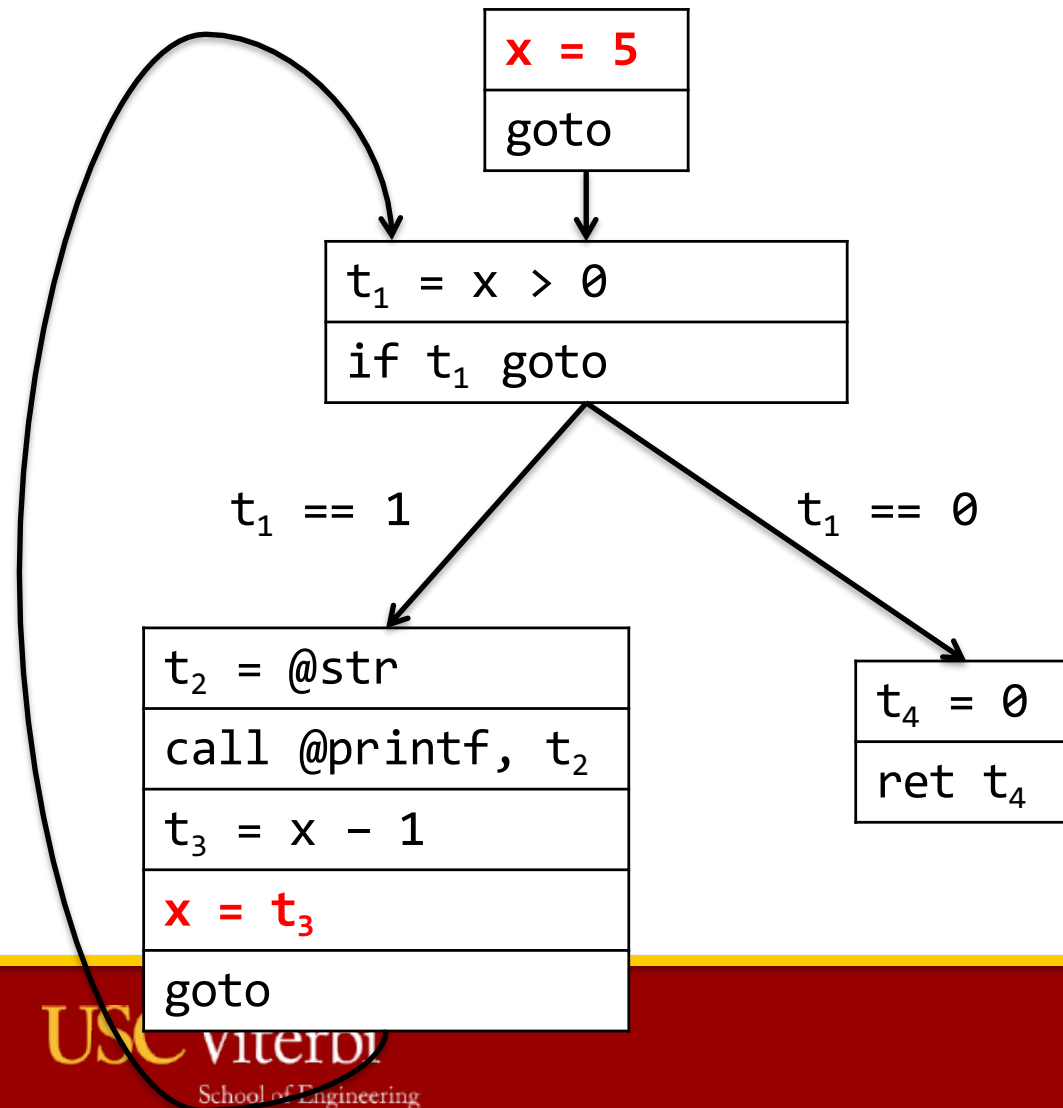


```
int main() {  
    int x = 5;  
    while (x > 0) {  
        printf("Hello!");  
        --x;  
    }  
    return 0;  
}
```

Static Single Assignment Form



- In static **single assignment form** (SSA) form, each variable in the IR can only have one assignment statement




```
int main() {  
    int x = 5;  
    while (x > 0) {  
        printf("Hello!");  
        --x;  
    }  
    return 0;  
}
```

NOT SSA FORM ☹️

A Simpler Example



- Still not SSA form ☹️

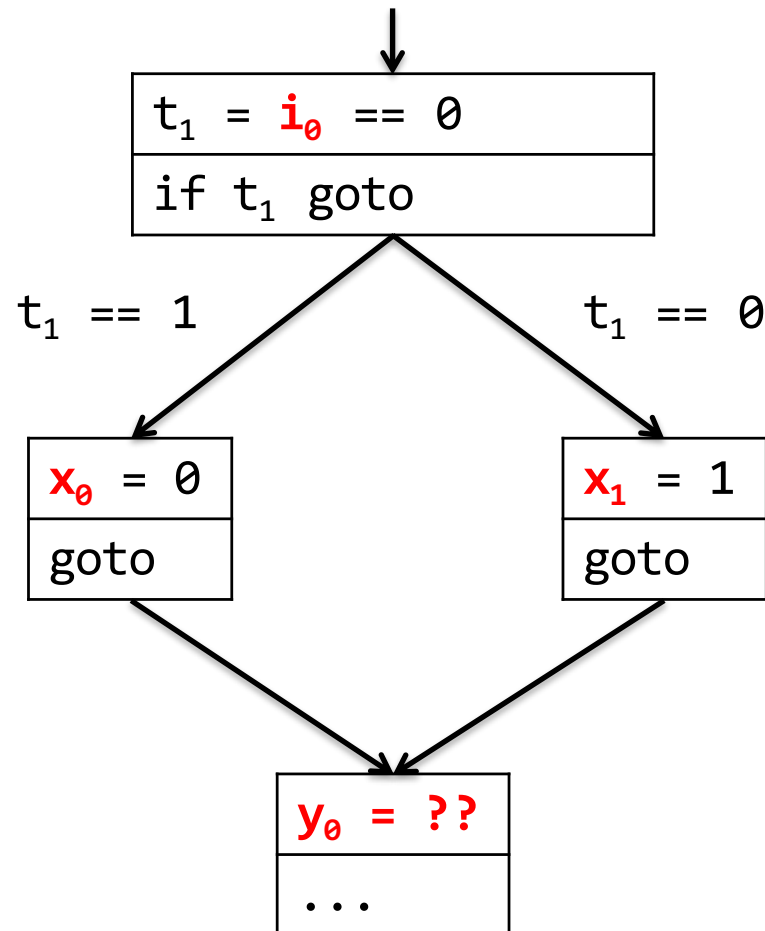


```
if (i == 0) {  
    x = 0;  
} else {  
    x = 1;  
}  
  
y = x;
```

Converting to SSA



- We can add a subscript to each variable assignment to make it unique

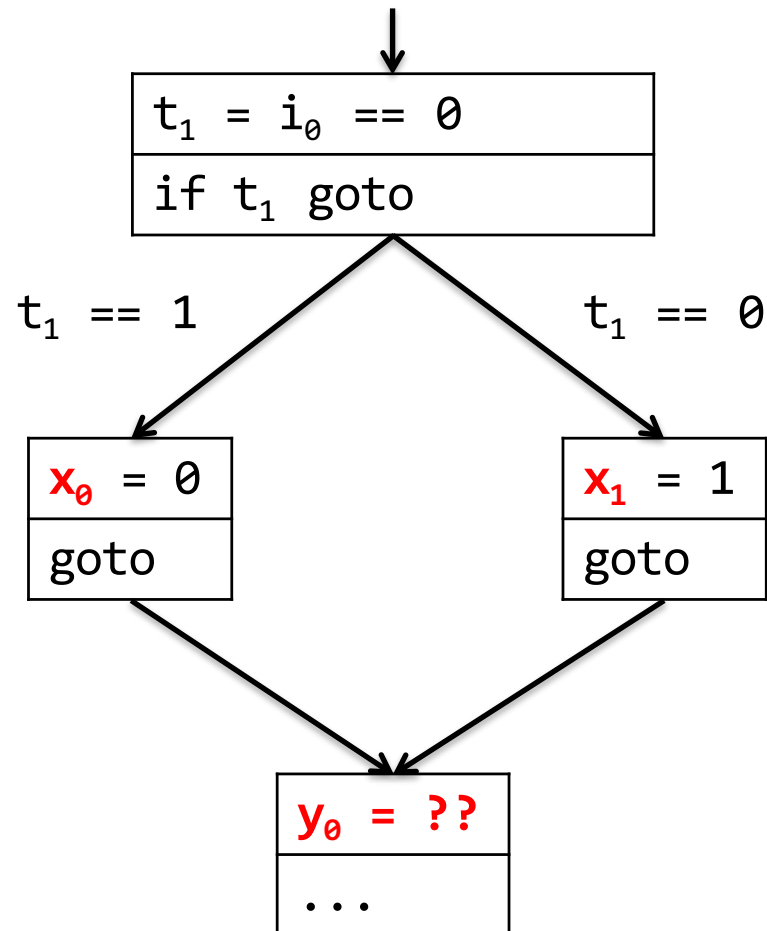


```
if (i == 0) {  
    x = 0;  
} else {  
    x = 1;  
}  
  
y = x;
```

Converting to SSA



- Problem:** Should y_0 be set to x_0 or x_1 ?

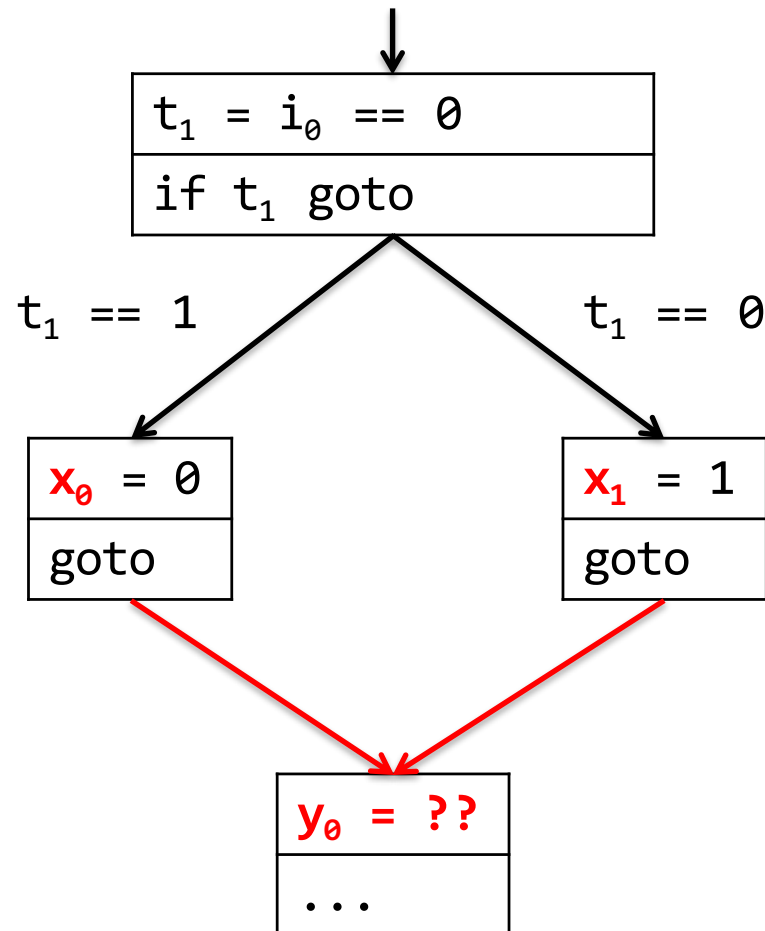


```
if (i == 0) {  
    x = 0;  
} else {  
    x = 1;  
}  
  
y = x;
```



Converting to SSA

- **Problem:** Should y_0 be set to x_0 or x_1 ?
- ***It depends on which basic block we came from!***

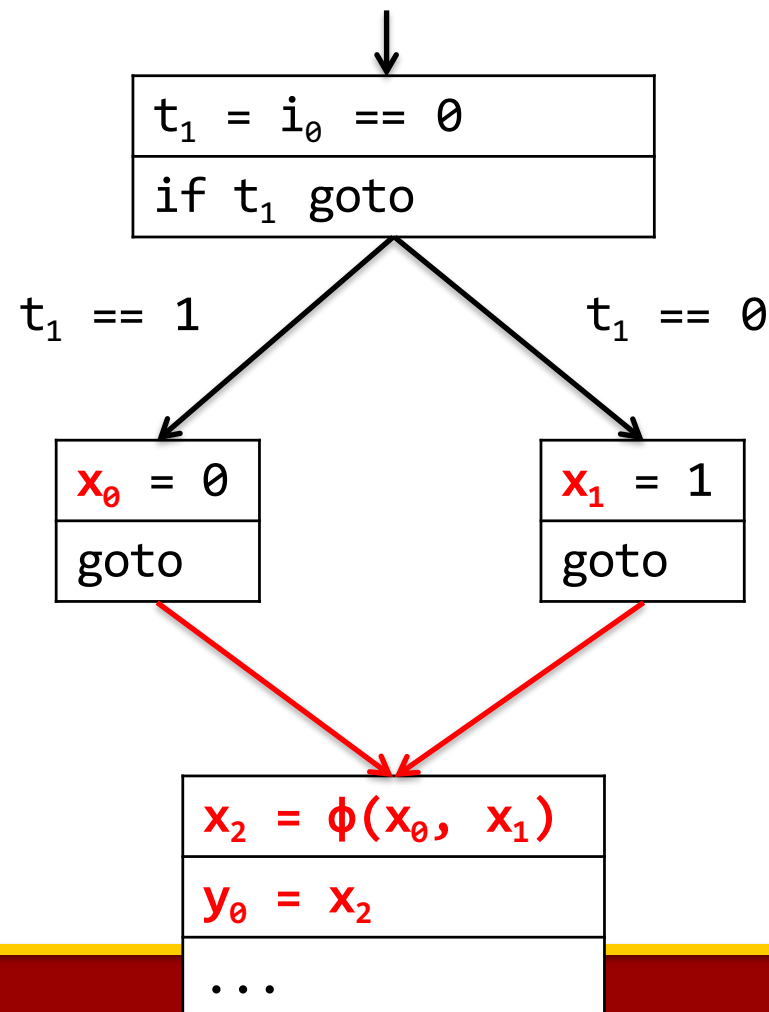


```
if (i == 0) {  
    x = 0;  
} else {  
    x = 1;  
}  
  
y = x;
```



Converting to SSA – Phi Nodes

- A **phi node** (or **ϕ -node**) is a special instruction that will select from multiple options based on the incoming edge in the CFG

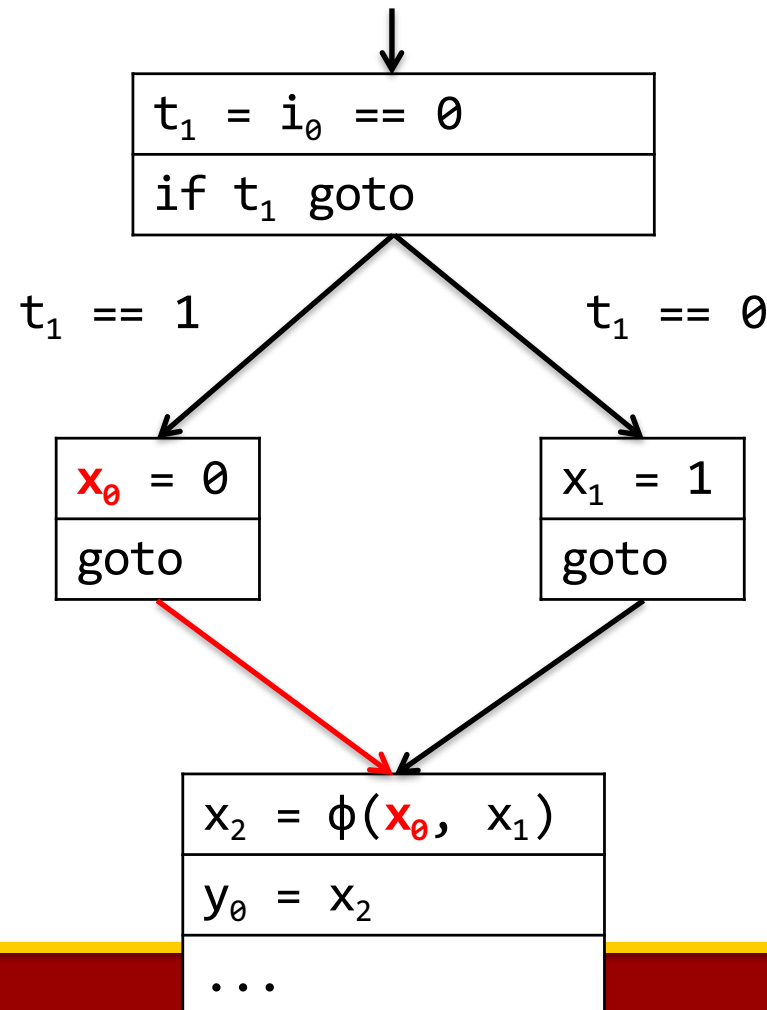


```
if (i == 0) {  
    x = 0;  
} else {  
    x = 1;  
}  
  
y = x;
```

Converting to SSA – Phi Nodes



- If control flow comes from the *left* predecessor, then $x_2 = x_0$

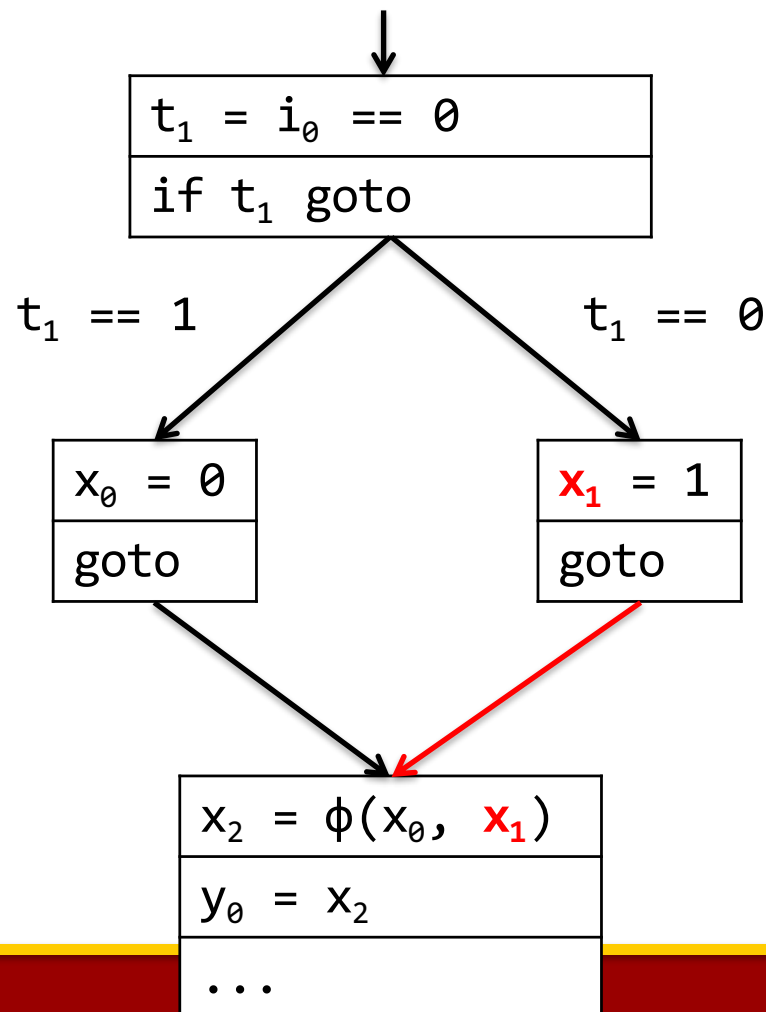


```
if (i == 0) {  
    x = 0;  
} else {  
    x = 1;  
}  
  
y = x;
```

Converting to SSA – Phi Nodes



- If control flow comes from the **right** predecessor, then $x_2 = x_1$



```
if (i == 0) {  
    x = 0;  
} else {  
    x = 1;  
}  
  
y = x;
```

Three Modes of the LLVM IR



- The LLVM IR can be used in three different ways:
 1. As a text file on disk, that looks a lot like assembly
 2. As a binary file on disk, which can be compiled to native code or interpreted
 3. As a set of data structures loaded in memory, used by the compiler

Three Modes of the LLVM IR



- First, let's focus on #1

- 1. As a text file on disk, that looks a lot like assembly**
2. As a binary file on disk, which can be compiled to native code or interpreted
3. As a set of data structures loaded in memory, used by the compiler

An Example



- Last time we had the following simple USC program:

```
int main() {  
    int x = 5;  
    while (x > 0) {  
        printf("Hello!");  
        --x;  
    }  
    return 0;  
}
```

- Let's look at it in LLVM IR...

In LLVM IR



```
; ModuleID = 'main'
@.str = private unnamed_addr constant [7 x i8] c"Hello!\00"
declare i32 @printf(i8*, ...)
define i32 @main() {
entry:
    br label %while.cond

while.cond:    ; preds = %while.body, %entry
    %x = phi i32 [ %dec, %while.body ], [ 5, %entry ]
    %cmp = icmp sgt i32 %x, 0
    %0 = zext i1 %cmp to i32
    %tobool = icmp ne i32 %0, 0
    br i1 %tobool, label %while.body, label %while.end

while.body:    ; preds = %while.cond
    %1 = i8* getelementptr inbounds ([7 x i8]* @.str, i32 0, i32 0)
    %2 = call i32 @printf(%1)
    %dec = sub i32 %x, 1
    br label %while.cond

while.end:    ; preds = %while.cond
    ret i32 0
}
```

Some Basic Properties...



```
; ModuleID = 'main'
@.str = private unnamed_addr constant [7 x i8] c"Hello!\00"
declare i32 @printf(i8*, ...)
define i32 @main() {
entry:
    br label %while.cond

while.cond:    ; preds = %while.body, %entry
    %x = phi i32 [ %dec, %while.body ], [ 5, %entry ]
    %cmp = icmp sgt i32 %x, 0
    %0 = zext i1 %cmp to i32
    %tobool = icmp ne i32 %0, 0
    br i1 %tobool, label %while.body, label %while.end

while.body:    ; preds = %while.cond
    %1 = i8* getelementptr inbounds ([7 x i8]* @.str, i32 0, i32 0)
    %2 = call i32 @printf(i8*, ...)* @printf(%1)
    %dec = sub i32 %x, 1
    br label %while.cond

while.end:    ; preds = %while.cond
    ret i32 0
}
```

Comments begin
w/ semi-colon

Some Basic Properties...



```
; ModuleID = 'main'
@.str = private unnamed_addr constant [7 x i8] c"Hello!\00"
declare i32 @printf(i8*, ...)
define i32 @main() {
entry:
    br label %while.cond

while.cond:    ; preds = %while.body, %entry
    %x = phi i32 [ %dec, %while.body ], [ 5, %entry ]
    %cmp = icmp sgt i32 %x, 0
    %0 = zext i1 %cmp to i32
    %tobool = icmp ne i32 %0, 0
    br i1 %tobool, label %while.body, label %while.end

while.body:    ; preds = %while.cond
    %1 = i8* getelementptr inbounds ([7 x i8]* @.str, i32 0, i32 0)
    %2 = call i32 @printf(i8*, ...)* @printf(%1)
    %dec = sub i32 %x, 1
    br label %while.cond

while.end:    ; preds = %while.cond
    ret i32 0
}
```

Function
implementations
are enclosed in
braces, just like
in C/C++

Some Basic Properties...



```
; ModuleID = 'main'
@.str = private unnamed_addr constant [7 x i8] c"Hello!\00"
declare i32 @printf(i8*, ...)
define i32 @main() {
entry:
    br label %while.cond

while.cond:    ; preds = %while.body, %entry
    %x = phi i32 [ %dec, %while.body ], [ 5, %entry ]
    %cmp = icmp sgt i32 %x, 0
    %0 = zext i1 %cmp to i32
    %tobool = icmp ne i32 %0, 0
    br i1 %tobool, label %while.body, label %while.end

while.body:    ; preds = %while.cond
    %1 = i8* getelementptr inbounds ([7 x i8]* @.str, i32 0, i32 0)
    %2 = call i32 @printf(i8*, ...)* @printf(%1)
    %dec = sub i32 %x, 1
    br label %while.cond

while.end:    ; preds = %while.cond
    ret i32 0
}
```

Each basic block begins with the name (label) of the block, followed by comments referencing any predecessor blocks.

Some Basic Properties...



```
; ModuleID = 'main'
@.str = private unnamed_addr constant [7 x i8] c"Hello!\00"
declare i32 @printf(i8*, ...)
define i32 @main() {
entry:
    br label %while.cond

while.cond:    ; preds = %while.body, %entry
    %x = phi i32 [ %dec, %while.body ], [ 5, %entry ]
    %cmp = icmp sgt i32 %x, 0
    %0 = zext i1 %cmp to i32
    %tobool = icmp ne i32 %0, 0
    br i1 %tobool, label %while.body, label %while.end

while.body:    ; preds = %while.cond
    %1 = i8* getelementptr inbounds ([7 x i8]* @.str, i32 0, i32 0)
    %2 = call i32 (i8*, ...)* @printf(%1)
    %dec = sub i32 %x, 1
    br label %while.cond

while.end:    ; preds = %while.cond
    ret i32 0
}
```

Variables or
virtual registers
are always
prefaced with a
% sign.

Basic block
names are
prefaced with %
only when used
as an operand

Some Basic Properties...



```
; ModuleID = 'main'
@.str = private unnamed_addr constant [7 x i8] c"Hello!\00"
declare i32 @printf(i8*, ...)
define i32 @main() {
entry:
    br label %while.cond

while.cond:    ; preds = %while.body, %entry
    %x = phi i32 [ %dec, %while.body ], [ 5, %entry ]
    %cmp = icmp sgt i32 %x, 0
    %0 = zext i1 %cmp to i32
    %tobool = icmp ne i32 %0, 0
    br i1 %tobool, label %while.body, label %while.end

while.body:    ; preds = %while.cond
    %1 = i8* getelementptr inbounds ([7 x i8]* @.str, i32 0, i32 0)
    %2 = call i32 (i8*, ...)* @printf(%1)
    %dec = sub i32 %x, 1
    br label %while.cond

while.end:    ; preds = %while.cond
    ret i32 0
}
```

The language is
strongly-typed

i32 = 32-bit int

i1 = bool

i8 = char

i8* = pointer to
char

7 x i8 = array of 7
characters

Some Basic Properties...



```
; ModuleID = 'main'
@.str = private unnamed_addr constant [7 x i8] c"Hello!\00"
declare i32 @printf(i8*, ...)
define i32 @main() {
entry:
    br label %while.cond

while.cond:    ; preds = %while.body, %entry
    %x = phi i32 [ %dec, %while.body ], [ 5, %entry ]
    %cmp = icmp sgt i32 %x, 0
    %0 = zext i1 %cmp to i32
    %tobool = icmp ne i32 %0, 0
    br i1 %tobool, label %while.body, label %while.end

while.body:    ; preds = %while.cond
    %1 = i8* getelementptr inbounds ([7 x i8]* @.str, i32 0, i32 0)
    %2 = call i32 (i8*, ...)* @printf(%1)
    %dec = sub i32 %x, 1
    br label %while.cond

while.end:    ; preds = %while.cond
    ret i32 0
}
```

It as phi nodes,
which means SSA
form.

But...

SSA Form and LLVM IR



- There are two ways data can be stored in LLVM IR:
 - In virtual registers, which must be in SSA form
 - On the stack/memory, which is **not** in SSA form (eg. you can write to the same memory address multiple times)
- For simplicity, in PA3 we'll use the stack for all declared variables, and virtual registers only for temporary computations
- We'll worry about SSA form in PA4

The Prior Example w/ the Stack...



```
; ModuleID = 'main'
@.str = private unnamed_addr constant [7 x i8] c"Hello!\00"
declare i32 @printf(i8*, ...)
define i32 @main() {
entry:
    %x = alloca i32
    store i32 5, i32* %x
    br label %while.cond

while.cond:      ; preds = %while.body, %entry
    %x1 = load i32* %x
    %cmp = icmp sgt i32 %x1, 0
    %0 = zext i1 %cmp to i32
    %tobool = icmp ne i32 %0, 0
    br i1 %tobool, label %while.body, label %while.end

while.body:      ; preds = %while.cond
    %1 = i8* getelementptr inbounds ([7 x i8]* @.str, i32 0, i32 0)
    %2 = call i32 @printf(i8*, ...)* @printf(%1)
    %x2 = load i32* %x
    %dec = sub i32 %x2, 1
    store i32 %dec, i32* %x
    br label %while.cond

while.end:      ; preds = %while.cond
    ret i32 0
}
```

LLVM Instructions of Note



- We'll only cover the ones you'll need to use in this class
- Most instructions have a lot of optional parameters, but I've pared it down to the bare minimum
- Full Language Reference: <http://llvm.org/docs/LangRef.html>

alloca Instruction



- Allocates memory on the stack and returns a pointer to the memory

- Syntax:

`<result> = alloca <type> [, <ty> <NumElements>]`

- Examples:

`%x = alloca i32 ; Allocates one 32-bit value`

`%y = alloca i32, i32 5 ; Allocates an array of 5 i32s`

store Instruction



- Stores a value into a memory address

- Syntax:

`store` <type> <value>, <ty>* <pointer>

- Example:

`%x = alloca i32 ; Allocates one 32-bit value`

`store i32 20, i32* %x ; Store 20 in the address`

load Instruction



- Read data from memory

- Syntax:

`<result> = load <ty>* <pointer>`

- Example:

`%x = alloca i32 ; Allocates one 32-bit value`

`store i32 20, i32* %x ; Store 20 in the address`

`%val = load i32* %x ; Loads *x (20) into val`



Binary Operators

- All of the binary operators follow essentially the same syntax:

`<result> = <instr> <ty> <op1>, <op2>`

- Operators of note:

- `add` – Integer addition
- `sub` – Integer subtraction
- `mul` – Integer multiplication
- `sdiv` – Signed integer division
- `srem` – Signed integer remainder/modulus

- Examples:

`%a = add i32 %x, %y ; %a = %x + %y`

`%b = sub i32 %a, 5 ; %b = %a - 5`

sext Instruction



- Signed extend from a smaller integral type to a larger integral type

- Syntax:

`<result> = sext <ty> <value> to <ty2>`

- Example:

`%x = sext i8 %a to i32 ; Extend from 8 to 32 bits`

trunc Instruction



- Truncate from a larger integral type to a smaller integral type

- Syntax:

`<result> = trunc <ty> <value> to <ty2>`

- Example:

`%x = trunc i32 %a to i8 ; Truncate from 32 to 8 bits`

icmp Instruction



- Compares two integral values (or pointers), and returns a bool result of the comparison

- Syntax:

`<result> = icmp <cond> <ty> <op1>, <op2>`

...where `<cond>` is the condition such as `eq`, `ne`, ...

- Examples:

`%b = icmp eq i32 %x, %y ; %b = %x == %y`

`%c = icmp sgt i32 %x, %y ; %c = %x > %y`



ret Instruction

- Returns from a function. This is a terminator instruction. The type returned must match the function signature

- Syntax:

`ret <type> <value>`

or

`ret void`

- Examples:

`ret i32 %x ; Returns the 32-bit integer %x`

`ret void ; Returns void`

br Instruction



- Branch – can be either conditional or unconditional. This is a terminator instruction.

- Syntax:

`br label <dest>`

or

`br i1 <cond>, label <iftrue>, label <iffalse>`

- Examples:

`br label %bb0 ; Unconditional jump to bb0`

`; (%x ? goto %bb0 : goto %bb1)`

`br i1 %x, label %bb0, label %bb1`

phi Instruction



- Used for SSA form phi nodes

- Syntax:

`<result> = phi <ty> [<val0>, <label0>], ...`

- Example:

`; %x = 20 if coming from %bb0, or %a if from %bb1`

`%x = phi i32 [20, %bb0], [%a, %bb1]`

getelementptr Instruction



- Used to get the address of an element in arrays and structs (among other things)
- So confusing that there even is an entire doc on it:
<http://llvm.org/docs/GetElementPtr.html>
- Don't worry about the syntax!



- LLVM also supports some slightly higher level intrinsic functions, such as some useful Standard C library functions:
 - `llvm.memcpy`
 - `llvm.memset`
 - `llvm.sqrt`
 - `llvm.pow`
 - `llvm.ceil/floor`

The Prior Example w/ the Stack...



```
; ModuleID = 'main'
@.str = private unnamed_addr constant [7 x i8] c"Hello!\00"
declare i32 @printf(i8*, ...)
define i32 @main() {
entry:
    %x = alloca i32
    store i32 5, i32* %x
    br label %while.cond

while.cond:      ; preds = %while.body, %entry
    %x1 = load i32* %x
    %cmp = icmp sgt i32 %x1, 0
    %0 = zext i1 %cmp to i32
    %tobool = icmp ne i32 %0, 0
    br i1 %tobool, label %while.body, label %while.end

while.body:      ; preds = %while.cond
    %1 = i8* getelementptr inbounds ([7 x i8]* @.str, i32 0, i32 0)
    %2 = call i32 @printf(i8*, ...)* @printf(%1)
    %x2 = load i32* %x
    %dec = sub i32 %x2, 1
    store i32 %dec, i32* %x
    br label %while.cond

while.end:      ; preds = %while.cond
    ret i32 0
}
```

```
int main() {
    int x = 5;
    while (x > 0) {
        printf("Hello!");
        --x;
    }
    return 0;
}
```

Three Modes of the LLVM IR



- Ok, what about some stuff on #3?
 1. As a text file on disk, that looks a lot like assembly
 2. As a binary file on disk, which can be compiled to native code or interpreted
 3. **As a set of data structures loaded in memory, used by the compiler**



- The Module class corresponds to all code in one object file
- Contains things such as:
 - List of all functions in the Module
 - List of all global variables
 - LLVM's internal SymbolTable (that you should pretty much never touch)



- *Most* of the types you'll be using inherit from `llvm::Value`
- One of the features `llvm::Value` provides is a custom RTTI implementation, via:

```
// isa returns true if value is-a pointer to Type
```

```
isa<Type>(value)
```

```
// Returns pointer to Type if value is-a pointer to Type
```

```
// Otherwise, returns nullptr
```

```
dyn_cast<Type>(value);
```

```
// Like dyn_cast, except it asserts if the cast fails
```

```
cast<Type>(value);
```



- Encapsulates a function
- Inherits (eventually) from llvm::Value
- Allows you to do things such as:
 - Get the entry block of the function
 - Iterate over all of the basic blocks in a function
 - Access/iterate over the arguments to the function



- Corresponds to a basic block
- Inherits (eventually) from llvm::Value
- Allows you to do things such as:
 - Iterate over all of the linear instructions in the basic block
 - Get the terminator instruction



- Every instruction inherits from this
- Since llvm::Instruction inherits from llvm::Value, every instruction can also be treated as a value
- This makes it really simple to pass the result of an instruction as an operand of another instruction



- [IRBuilder](#) is absolutely your best friend when generating LLVM IR
- It has a factory method to create every instruction, which prevents you from having to write out the IR in text form

- Example (from some of the provided code):

```
{  
    IRBuilder<> build(ctx.mBlock);  
    // We can assume it WILL be an i32 here  
    // since it'd have been zero-extended otherwise  
    lhsVal = build.CreateICmpNE(lhsVal, ctx.mZero, "tobool");  
    build.CreateCondBr(lhsVal, rhsBlock, endBlock);  
}
```