

Introduction to LLVM

Eric Christopher (@echristo) and Johannes Doerfert (jdoerfert@anl.gov)

Why are you here?

- Brief overview of LLVM
 - Understand some of the various pieces
 - Want to know where to go for more
-
- Everything in this talk easily could get its own 45 minute tutorial!

Source Layout

<code>llvm-project</code>	LLVM mono-repo (https://github.com/llvm/llvm-project)
<code>llvm-project/clang</code>	C/C++/Cuda/OpenCL/... language front-end
<code>llvm-project/llvm</code>	"core" LLVM (middle-end & backends)
<code>llvm-project/{lld,lldb,...}</code>	"LLVM-based" tools, linker, debugger, ...

all subprojects look similar: `include`, `lib`, `tests`, `utils`

(see also <http://llvm.org/docs/GettingStarted.html#general-layout>)

Configure & Build

Single command often suffices to configure:

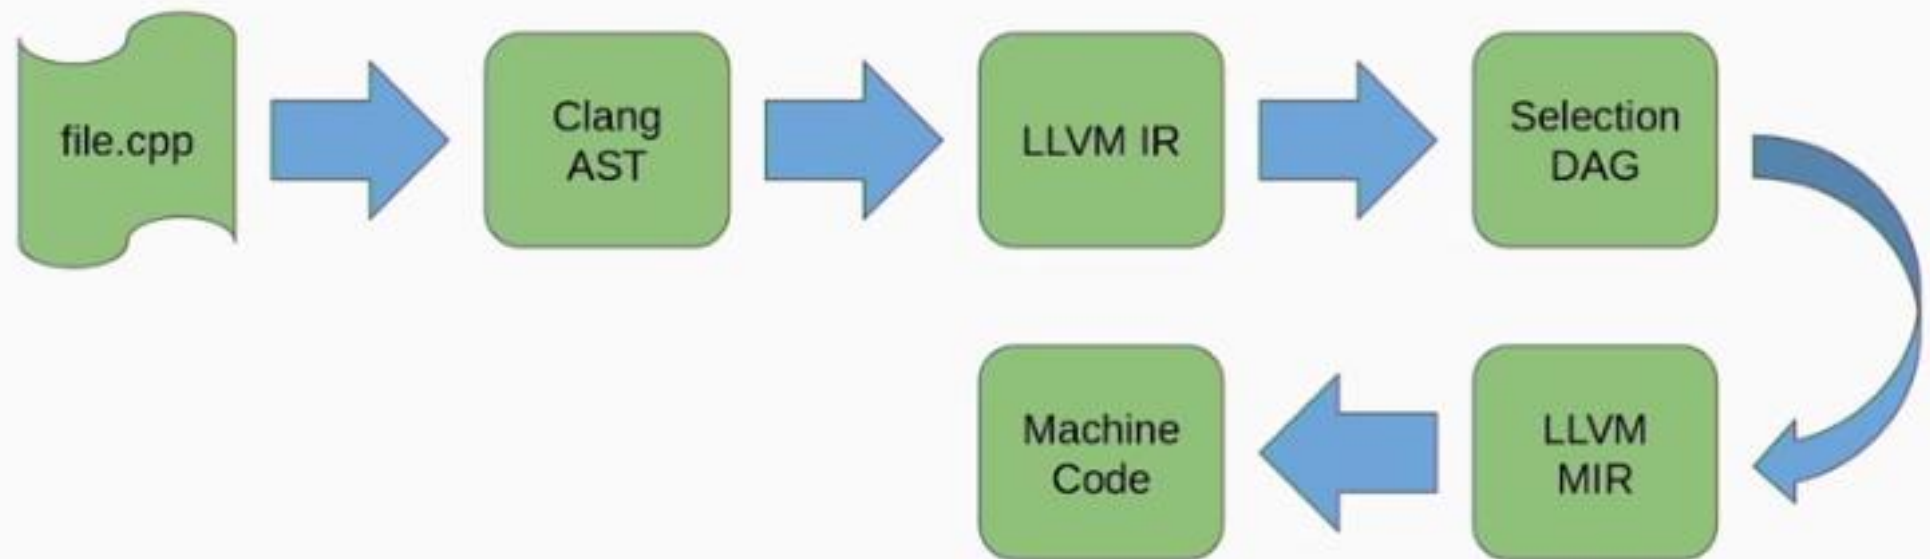
```
cmake /src/llvm-project/llvm -DLLVM_ENABLE_PROJECTS='clang;openmp'  
make -j
```

Useful options include: CMAKE_BUILD_TYPE={Release,Asserts,...}
LLVM_ENABLE_ASSERTIONS={ON,OFF}
LLVM_CCACHE_BUILD={ON,OFF}
-G Ninja


Various resources available online! Start here:

<http://llvm.org/docs/GettingStarted.html>

Compilers 101




Compilers 101



Life of a Program: Control Flow Graph (CFG)

```
long f(long a, long b) {  
    long x = a;  
    if (a > b)  
        x += 20;  
    else  
        x += b;  
    return x;  
}
```



Life of a Program: C -> LLVM-IR

```
clang -O3 -Xclang -disable-llvm-passes           \ # prepare for O3 but don't run it
    -S -emit-llvm code.c -o code.ll              \ # emit LLVM IR for code.c into code.ll
opt -S -mem2reg -instnamer code.ll -o code_before_opt.ll \ # slight cleanup
```



```
; ModuleID = 'f.ll'
source_filename = "f.c"
target datalayout = "e-m:e-p270:32:32-p271:32:32-p272:64:64-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"
```

```
; Function Attrs: nounwind uwtable
define dso_local i64 @f(i64 %a, i64 %b) #0 {
```

```
if.then:                                     ; preds = %entry
    %add = add nsw i64 %a, 20
    br label %if.end
```

```
if.else:                                     ; preds = %entry
    %addl = add nsw i64 %a, %b
    br label %if.end
```

```
}
```


```
attributes #0 = { nounwind uwtable }
```

```
!llvm.module.flags = !{!0}
```

```
!llvm.ident = !{!1}
```

```
!0 = !{i32 1, !"wchar_size", i32 4}
```

```
!1 = !{!"clang version 10.0.0"}
```



```
; ModuleID = 'f.ll'
source_filename = "f.c"
target datalayout = "e-m:e-p270:32:32-p271:32:32-p272:64:64-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"
```

```
; Function Attrs: nounwind uwtable
```

```
define dso local i64 @f(i64 %a, i64 %b) #0 {
```

```
entry:
```

```
    %cmp = icmp sgt i64 %a, %b
    br i1 %cmp, label %if.then, label %if.else
```

```
if.then:
```

```
    %add = add nsw i64 %a, 20
    br label %if.end
```

```
; preds = %entry
```

```
if.else:
```

```
    %add1 = add nsw i64 %a, %b
    br label %if.end
```

```
; preds = %entry
```

```
if.end:
```

```
    %x.0 = phi i64 [ %add, %if.then ], [ %add1, %if.else ]
    ret i64 %x.0
```

```
; preds = %if.else, %if.then
```


```
attributes #0 = { nounwind uwtable }
```

```
!llvm.module.flags = !{!0}
```

```
!llvm.ident = !{!1}
```

```
!0 = !{i32 1, !"wchar_size", i32 4}
```

```
!1 = !{!"clang version 10.0.0"}
```



LLVM-IR Hierarchy

- `llvm::Module`
 - `llvm::GlobalVariable`
 - `llvm::Function` (declarations & definitions)
 - `llvm::BasicBlock`
 - `llvm::Instruction`
 - `llvm::ICmpInst`
 - `llvm::BranchInst`

LLVM Intermediate Representation (LLVM-IR)

```
%sum = add i32 4, %var           ; Binary operations
%cmp = icmp sge i32 %a, %b

%value = load i32, i32* %ptr      ; Memory operations
store i32 %value, i32* %ptr2

br label %next-block             ; Terminator operations
br i1 %cmp, label %then-block, label %else-block
ret float %a
ret void

%X = trunc i32 257 to i8         ; Cast operations
%Y = sext i32 %V to i64
%Z = bitcast i8* %q to i32*

%ptr = alloca [4 x double]       ; Other operations
%ret = call i32 @foo(i8* %fmt, i16 %val)
%phi = phi float [ %valA, %predBlockA ], [ %valB, %predBlockB ]
%elemI = getelementptr float, float* %A, i64 %I
```

LLVM Intermediate Representation (LLVM-IR)

assembly-like language **always** in SSA-form with infinite registers.

Organization:

<code>module</code>	= list of global symbols (functions, global variables, ...)
<code>function</code>	= list of basic blocks that form a control-flow-graph (CFG)
<code>basic block</code>	= list of instructions, terminated by a branch/return/...
<code>instruction</code>	= typed assembly instruction with operands
<code>constant</code>	= constant literals, globals, ...
<code>value</code>	= almost anything, mostly constants and instructions

LLVM Intermediate Representation (LLVM-IR)

assembly-like language **always** in SSA-form with infinite registers.

- “global symbols” start with “@”
- “local symbols” start with “%”
- Basic block names start with “%” when used
- Basic block names end with “:” when defined

LLVM Intermediate Representation (LLVM-IR)

Common pitfalls:

- control-flow can be irreducible and irreducible loops are “not recognized”
- only *reachable* LLVM-IR has to be *well-formed*, e.g., wrt. SSA properties
- address spaces are part of the pointer type (and not to be ignored!)
- `null` (= 0) can be a valid pointer, even in address space 0 (=default)
- types are sign agnostic
- `llvm::Constants` include values “constant at startup time”, e.g., addr of globals

Useful command line options

`clang -mllvm <option for opt>`

`--help & --help-hidden`

`--debug & --debug-only=attributor,must-execute`

`--debug-pass={Arguments,Structure,Executions,Details}`

`--print-before-all/--print-after-all/--print-XXXX/--dot-XXXX`

Targets and Code Generation

Codegen

- Instruction Selectors
- Register Allocation
- Instruction Scheduling
- Machine IR
- Machine Optimizations
- Machine Specific Features (stack protection, etc)
- Exception Handling
- "Assembly Printing"
- Target Dependency Breaking
- Calls heavily into Target code

Target

- x86, ARM, AArch64, WebAssembly, Power, Hexagon, CPUs, GPUs, Bytecodes
- Assembly Parsing/Disassembly
- Instruction Descriptions
- Optimization Passes
- CPU pipeline descriptions
- Calling Conventions and ABI
- Machine Code and Object File

Targets and Code Generation - An Example

```
def LEA64_32r : I<0x8D, MRMSrcMem,  
    (outs GR32:$dst), (ins lea64_32mem:$src),  
    "lea{l}\t{$src|$dst}, {$dst|$src}",  
    [(set GR32:$dst, lea64_32addr:$src)]>,  
    OpSize32, Requires<[In64BitMode]>;
```

Targets and Subtargets

Target/

- TargetMachine contains Module level lowering, data, and transform information
- Module level means object file/OS ABI
- Subtarget contains Function level lowering information.
- Primarily initialized with CPU and Features
- SubtargetFeatures include things like ISA level

Requires<[In64BitMode]>

X86.td:

```
def In64BitMode :  
  Predicate<"Subtarget->is64Bit()">
```

X86Subtarget.h,.cpp

```
bool is64Bit() { return In64BitMode;}
```

```
In64BitMode(TargetTriple.getArch() ==  
Triple::x86_64)
```

Machine IR

CodeGen/MachineInstr.h

Target/

- Follows a similar hierarchy to the rest of LLVM IR
- Dependent upon LLVM IR constructs for
- Extensive APIs to construct MIR
- Contains target-dependent and target-independent opcodes, registers
- Can be produced out of the llc program

X86FastISel.cpp:

```
unsigned Opc = TLI.getPointerTy(DL) ==  
MVT::i32 ?  
(Subtarget->isTarget64BitILP32() ?  
X86::LEA64_32r : X86::LEA32r)  
: X86::LEA64r;  
  
addFullAddress(BuildMI(*FuncInfo.MBB,  
FuncInfo.InsertPt, DbgLoc,  
TII.get(Opc), ResultReg), AM);
```

Instruction Selection and Register Allocation

CodeGen/

- Instruction pattern describes what is happening.
- 3 Instruction Selectors
 - SelectionDAG
 - FastISel
 - GlobalISel
- Operands describe what and where
- 2* Register Allocators
 - RegAllocFast
 - RegAllocGreedy

SelectionDAG View

```
(set GR32:$dst,  
lea64_32addr:$src)
```

```
X86InstrInfo.td
```

```
def lea64_32addr : ComplexPattern<i32,  
5, "selectLEA64_32Addr"...
```

```
X86ISelDAGToDAG.cpp:
```

```
bool
```

```
X86DAGToDAGISel::selectLEA64_32Addr
```

Object Generation - Assembly “Printing”

CodeGen/AsmPrinter

LEA64_32r : l<0x8D

- Target independent MIR to Assembly translation
- Calls into backends

leal 485498096(%edx,%eax,4), %eax

X86AsmPrinter::PrintLeaMemReference

Target/*AsmPrinter

- Handles target specific opcode and operand printing

Object Generation - Assembly “Parsing”

AsmParser/

encoding: [0x8d,0x84,0x82,0xf0,0x1c,0xf0,0x1c]

MC/

- Handles target-independent object file construction and encoding

Target/AsmParser

Target/MCTargetDesc

- Parses target specific assembly and encodings
- Handles target-dependent details

Object Reading

Object/

DebugInfo/

- Parses object files and target-independent data
- Calls into the target for instruction decoding

Target/Disassembler

- Handles target specific instruction decoding

llvm-objdump.cpp:

```
Binary &Binary = *OBinary.getBinary()
```

```
ObjectFile *O = dyn_cast<ObjectFile>(&Binary)
```

```
DIContext DICtx = DWARFContext::create(*O)
```

```
DICtx->dump(outs(), DumpOpts)
```


More Tools

tools/

- Examples of APIs
 - Canonical usage
 - Updated frequently
-
- Can be moved into a library if we have a need for reuse